

SoPa++: Leveraging explainability from hybridized RNN, CNN and weighted finite-state neural architectures

Atreya Shankar, shankar@uni-potsdam.de
Cognitive Systems: Language, Learning, and Reasoning (M.Sc.)

Thesis Progress Update
University of Potsdam, WiSe 20/21

March 15, 2021

Recap

- 1 Source code refactored to work with latest PyTorch and CUDA versions
- 2 Facebook NLU multi-class intent detection data set from [Schuster et al. \(2019\)](#)
- 3 Performance is important for explainability, because it does not make sense to explain a poorly performing model
- 4 Focus would be to develop SoPa++ as a neural model which can be easily decomposed into a *globally* explainable model

Semirings

Kuich and Salomaa 1986:

A semiring is a set \mathbb{K} along with two binary associative operations \oplus (addition) and \otimes (multiplication) and two identity elements: $\bar{0}$ for addition and $\bar{1}$ for multiplication. Semirings require that addition is commutative, multiplication distributes over addition, and that multiplication by $\bar{0}$ annihilates, i.e., $\bar{0} \otimes a = a \otimes \bar{0} = \bar{0}$

- Generic semiring notation: $\langle \mathbb{K}, \oplus, \otimes, \bar{0}, \bar{1} \rangle$
- Max-sum semiring: $\langle \mathbb{R} \cup \{-\infty\}, \max, +, -\infty, 0 \rangle$
- Max-product semiring: $\langle \mathbb{R} \cup \{-\infty\}, \max, \times, -\infty, 1 \rangle$
- Many semirings are possible, but we stick to *max* based semirings

Weighted Finite-State Automata

$$p(\mathbf{x}) = \boldsymbol{\pi} \otimes \left(\bigotimes_{i=1}^n \mathbf{T}(x_i) \right) \otimes \boldsymbol{\eta} \quad (1)$$

$$s(\mathbf{z}) = \bigoplus_{\mathbf{x} \in \Pi(\mathbf{z})} p(\mathbf{x}) \quad (2)$$

- Path $\mathbf{x} = \langle x_1, x_2, \dots, x_n \rangle$ derives some string $\mathbf{z} = \langle z_1, z_2, \dots, z_m \rangle$
- $\Pi(\mathbf{z})$ is the set of all possible paths that derive string \mathbf{z}
- $\boldsymbol{\pi}, \boldsymbol{\eta}$ are start and end state vectors
- \mathbf{T} is a transition matrix defining transition scores between states
- $p(\mathbf{x})$ refers to a score through a particular path
- $s(\mathbf{z})$ refers to an aggregate score through all possible paths
- Utilize Viterbi algorithm with *max* based semirings to compute $s(\mathbf{z})$
(Viterbi, 1967)

Finite transition types




Figure 1: Schematic for finite-state transitions; ϵ represents an epsilon-transition, “funny” and “magical” form self-loops and all other transitions are main transitions (Schwartz, Thomson, and Smith, 2018)

- Vanilla SoPa allowed 3 transitions types: main, epsilon and self-loops
- **SoPa++ uses only two transitions: namely main and wildcard transitions**
- Both consume one token and one state, which leads to pattern-string determinism and could make the *explainability* step easier

SoPa++: Lower Model

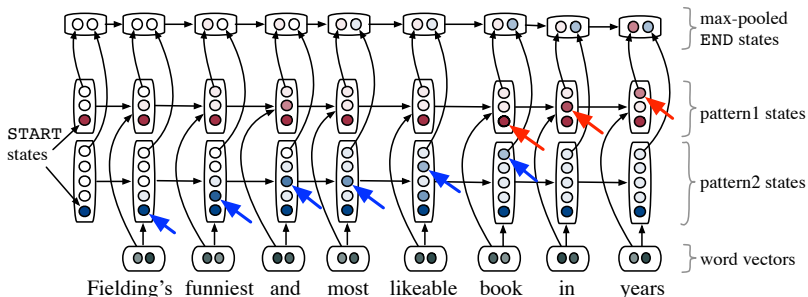


Figure 2: SoPa++ lower model schematic where pattern 1 is activated by the substring "in years" and pattern 2 is activated by the substring "funniest and most likeable book" (Schwartz, Thomson, and Smith, 2018)

- Vanilla SoPa places a MLP on top of the max-pooled states to retrieve classification results
- This contributes to a lack of explainability since MLPs are generally not easily explainable

SoPa++: Upper Model

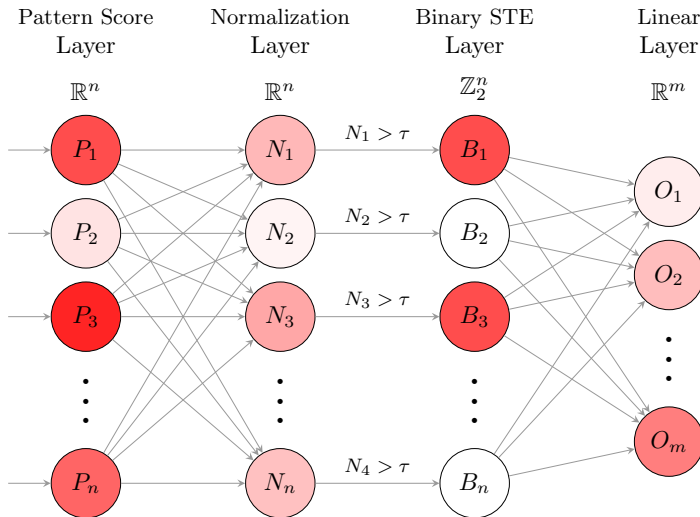


Figure 3: SoPa++ upper model schematic (combined lower-upper schematic under-construction); STE refers to a straight-through-estimator (Yin et al., 2019)

SoPa++: Explainable Model

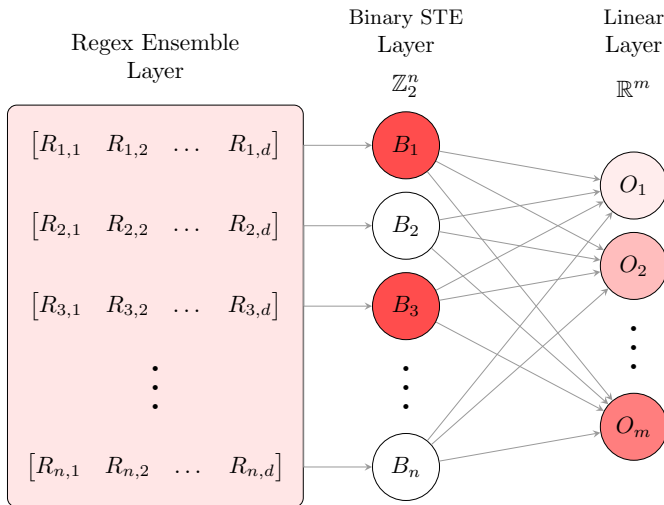


Figure 4: SoPa++ explainable variant derived from antecedent neural model

Training and Convergence




Figure 5: Excerpt from TensorBoard training logs for SoPa++ grid-search

- Model training procedure tends to converge fast
- Best model checkpoint(s) were reached within 1-6 epochs

Evaluation

Patterns	Test Precision	Test Recall	Test F_1
6-10_5-10_4-10_3-10	0.982	0.981	0.982
6-50_5-50_4-50_3-50	0.988	0.987	0.987

Table 1: Performance summary of various SoPa++ models; evaluation metrics are based on weighted averages

- Performance is competitive with that from literature and online benchmarks, which show F_1 scores ranging from 96-99% (Schuster et al., 2019)
- In both the above cases, the max-sum semiring outperforms the max-product semiring
- Even though the SoPa++ model variant with higher number of patterns performs better, a **smaller model** might be more beneficial for explainability

Explainability: Peek into the Regex Ensemble

Pattern length	Sample regular expressions
3	{* forecast}, {* activate}, {* weather}
4	{activate morning *}, {my alarm *}, {* minute snooze}
5	{add * reminder for}, {the boiled eggs *}
6	{to wish * * happy}, {* add * * to}

Table 2: Sample regular expressions that activate patterns of specified lengths; * is used for aesthetic purposes to represent a wildcard instead of the appropriate `\w+` regular expression

- Patterns above were processed from a small subset of the training data; the full processing is still under development
- Patterns are indicative of important (generalized) n-grams in text
- Analyzing the weights of the linear layer can tell us how important each pattern is for the prediction of each class

Next steps

- 1 Find a clever way for clustering/merging regular expressions to keep the size of explainable model small
- 2 Analyze the performance of both SoPa++ and its explainable variant on the test set to measure how different their performances are
- 3 For cases where the explainable model does not perform similar to the SoPa++ neural model, propose local sample explanations as alternative

Bibliography I

Kuich, Werner and Arto Salomaa (1986). “Linear Algebra”. In: *Semirings, automata, languages*. Springer, pp. 5–103.

Schuster, Sebastian et al. (June 2019). “Cross-lingual Transfer Learning for Multilingual Task Oriented Dialog”. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Minneapolis, Minnesota: Association for Computational Linguistics, pp. 3795–3805. DOI: 10.18653/v1/N19-1380. URL: <https://www.aclweb.org/anthology/N19-1380>.

Bibliography II

Schwartz, Roy, Sam Thomson, and Noah A. Smith (July 2018).

“Bridging CNNs, RNNs, and Weighted Finite-State Machines”. In: *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Melbourne, Australia: Association for Computational Linguistics, pp. 295–305.

DOI: 10.18653/v1/P18-1028. URL:

<https://www.aclweb.org/anthology/P18-1028>.

Viterbi, Andrew (1967). “Error bounds for convolutional codes and an asymptotically optimum decoding algorithm”. In: *IEEE transactions on Information Theory* 13.2, pp. 260–269.

Yin, Penghang et al. (2019). “Understanding straight-through estimator in training activation quantized neural nets”. In: *arXiv preprint arXiv:1903.05662*.