SoPa++: Leveraging explainability from hybridized RNN, CNN and weighted finite-state neural architectures

Atreya Shankar, `shankar@uni-potsdam.de`

Cognitive Systems: Language, Learning, and Reasoning (M.Sc.)

Thesis Progress Update

University of Potsdam, WiSe 20/21

February 3, 2021

Recap

1. Source code refactored to work with latest PyTorch and CUDA versions

2. Facebook NLU multi-class intent detection data set from Schuster et al. (2018)

3. Performance is important for explainability, because it does not make sense to explain a poorly performing model

4. Focus would be to develop SoPa++ as a neural model which can be easily decomposed into a *globally* explainable model

Semirings

Kuich and Salomaa 1986:
A semiring is a set $\mathbb{K}$ along with two binary associative operations $\oplus$ (addition) and $\otimes$ (multiplication) and two identity elements: $\bar{0}$ for addition and $\bar{1}$ for multiplication. Semirings require that addition is commutative, multiplication distributes over addition, and that multiplication by $\bar{0}$ annihilates, i.e., $\bar{0} \otimes a = a \otimes \bar{0} = \bar{0}$

- Generic semiring notation: $\langle \mathbb{K}, \oplus, \otimes, \bar{0}, \bar{1} \rangle$
- Max-sum semiring: $\langle \mathbb{R} \cup \{-\infty\}, \max, +, -\infty, 0 \rangle$
- Max-product semiring: $\langle \mathbb{R} \cup \{-\infty\}, \max, \times, -\infty, 1 \rangle$
- Many semirings are possible, but we stick to *max* based semirings

Weighted Finite-State Automata

$$p(\boldsymbol{x}) = \boldsymbol{\pi} \otimes \left(\bigotimes_{i=1}^{n} \boldsymbol{T}(x_i) \right) \otimes \boldsymbol{\eta} \qquad (1)$$

$$s(\boldsymbol{z}) = \bigoplus_{\boldsymbol{x} \in \Pi(\boldsymbol{z})} p(\boldsymbol{x}) \qquad (2)$$

- Path $\boldsymbol{x} = \langle x_1, x_2, \ldots, x_n \rangle$ derives some string $\boldsymbol{z} = \langle z_1, z_2, \ldots, z_m \rangle$
- $\Pi(\boldsymbol{z})$ is the set of all possible paths that derive string $\boldsymbol{z}$
- $\boldsymbol{\pi}$, $\boldsymbol{\eta}$ are start and end state vectors
- $\boldsymbol{T}$ is a transition matrix defining transition scores between states
- $p(\boldsymbol{x})$ refers to a score through a particular path
- $s(\boldsymbol{z})$ refers to an aggregate score through all possible paths
- Utilize Viterbi algorithm with *max* based semirings to compute $s(\boldsymbol{z})$
 (Viterbi, 1967)

4

Finite transition types

Figure 1: Schematic for finite-state transitions; ϵ represents an epsilon-transition, "funny" and "magical" form self-loops and all other transitions are main transitions (Schwartz, Thomson, and Smith, 2018)

- Vanilla SoPa allowed 3 transitions types: main, epsilon and self-loops
- **SoPa++ uses only two transitions: namely main and wildcard transitions**
- Both consume one token and one state, which leads to pattern-string determinism and could make the *explainability* step easier

SoPa++: Lower Model

Figure 2: SoPa++ lower model schematic where pattern 1 is activated by the substring "in years" and pattern 2 is activated by the substring "funniest and most likeable book" (Schwartz, Thomson, and Smith, 2018)

- Vanilla SoPa places a MLP on top of the max-pooled states to retrieve classification results

- This contributes to a lack of explainability since MLPs are generally not easily explainable
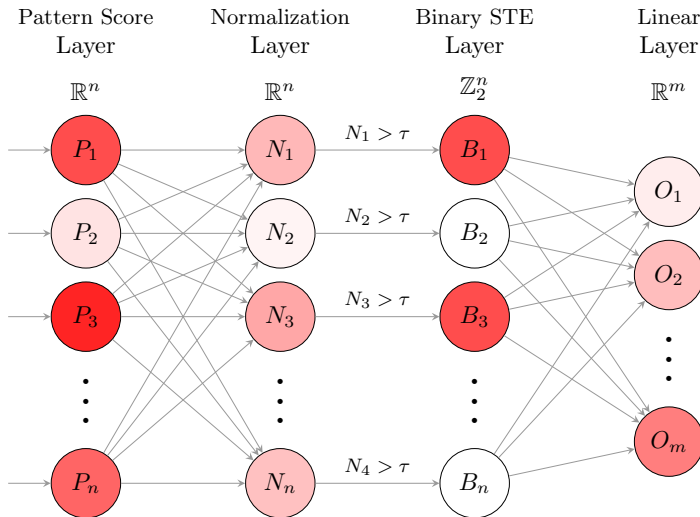
SoPa++: Upper Model

Figure 3: SoPa++ upper model schematic (combined lower-upper schematic under-construction); STE refers to a straight-through-estimator (Yin et al., 2019)
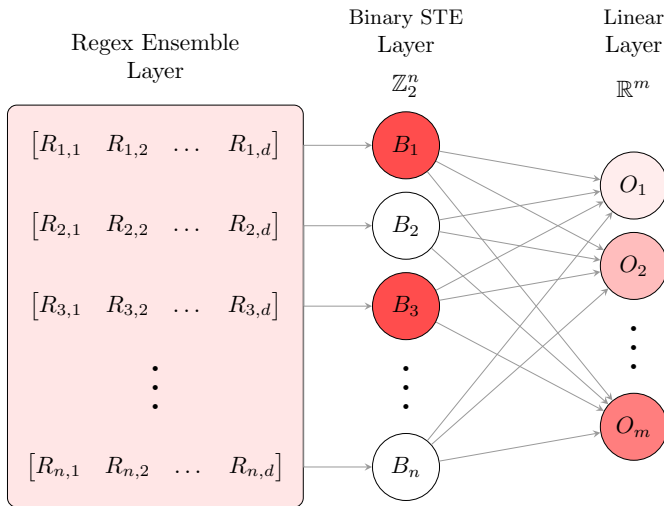
SoPa++: Explainable Model

Regex Ensemble Layer

Binary STE Layer $\mathbb{Z}_2^n$

Linear Layer $\mathbb{R}^m$

$$\begin{bmatrix} R_{1,1} & R_{1,2} & \ldots & R_{1,d} \end{bmatrix}$$

$$\begin{bmatrix} R_{2,1} & R_{2,2} & \ldots & R_{2,d} \end{bmatrix}$$

$$\begin{bmatrix} R_{3,1} & R_{3,2} & \ldots & R_{3,d} \end{bmatrix}$$

$$\vdots$$

$$\begin{bmatrix} R_{n,1} & R_{n,2} & \ldots & R_{n,d} \end{bmatrix}$$

B_1 B_2 B_3 $\vdots$ B_n

O_1 O_2 $\vdots$ O_m

Figure 4: SoPa++ explainable variant derived from antecedent neural model

Training and Convergence

Figure 5: Excerpt from TensorBoard training logs for SoPa++ grid-search

- Model training procedure tends to converge fast
- Best model checkpoint(s) were reached within 1-6 epochs

Evaluation

Patterns	Test Precision	Test Recall	Test F_1
6−10_5−10_4−10_3−10	0.982	0.981	0.982
6−50_5−50_4−50_3−50	0.988	0.987	**0.987**

Table 1: Performance summary of various SoPa++ models; evaluation metrics are based on weighted averages

- Performance is competitive with that from literature and online benchmarks, which show F_1 scores ranging from 96-99% (Schuster et al., 2018)
- In both the above cases, the max-sum semiring outperforms the max-product semiring
- Even though the SoPa++ model variant with higher number of patterns performs better, a **smaller model** might be more beneficial for explainability

Explainability: Peek into the Regex Ensemble

Pattern length	Sample regular expressions
3	{* forecast}, {* activate}, {* weather}
4	{activate morning *}, {my alarm *}, {* minute snooze}
5	{add * reminder for}, {the boiled eggs *}
6	{to wish * * happy}, {* add * * to}

Table 2: Sample regular expressions that activate patterns of specified lengths; * is used for aesthetic purposes to represent a wildcard instead of the appropriate \w+ regular expression

- Patterns above were processed from a small subset of the training data; the full processing is still under development
- Patterns are indicative of important (generalized) n-grams in text
- Analyzing the weights of the linear layer can tell us how important each pattern is for the prediction of each class

Next steps

1. Speed up regular expression extraction using vectorization; currently the process takes $\sim$1.5 hrs with a single-threaded process

2. Find a clever way for clustering/merging regular expressions to keep the size of explainable model small

3. Analyze the performance of both SoPa++ and its explainable variant on the test set to measure how different their performances are

4. For cases where the explainable model does not perform similar to the SoPa++ neural model, propose local sample explanations as alternative

Bibliography I

Arrieta, Alejandro Barredo et al. (2020). "Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI". In: *Information Fusion* 58, pp. 82–115.

Bocklisch, Tom et al. (2017). "Rasa: Open source language understanding and dialogue management". In: *arXiv preprint arXiv:1712.05181.*

Cybenko, George (1989). "Approximation by superpositions of a sigmoidal function". In: *Mathematics of control, signals and systems* 2.4, pp. 303–314.

Doran, Derek, Sarah Schulz, and Tarek R. Besold (2017). "What Does Explainable AI Really Mean? A New Conceptualization of Perspectives". In: *CoRR* abs/1710.00794. arXiv: 1710.00794. URL: http://arxiv.org/abs/1710.00794.

Bibliography II

Evans, Richard and Edward Grefenstette (2018). "Learning explanatory rules from noisy data". In: *Journal of Artificial Intelligence Research* 61, pp. 1–64.

Hornik, Kurt, Maxwell Stinchcombe, Halbert White, et al. (1989). "Multilayer feedforward networks are universal approximators.". In: *Neural networks* 2.5, pp. 359–366.

Hou, Bo-Jian and Zhi-Hua Zhou (2018). "Learning with Interpretable Structure from RNN". In: *CoRR* abs/1810.10708. arXiv: 1810.10708. URL: http://arxiv.org/abs/1810.10708.

Jiang, Chengyue et al. (2020). "Cold-Start and Interpretability: Turning Regular Expressions into Trainable Recurrent Neural Networks". In.

Kepner, Jeremy et al. (2018). "Sparse deep neural network exact solutions". In: *2018 IEEE High Performance extreme Computing Conference (HPEC)*. IEEE, pp. 1–8.

Bibliography III

Kuich, Werner and Arto Salomaa (1986). "Linear Algebra". In: *Semirings, automata, languages*. Springer, pp. 5–103.

Law, Mark, Alessandra Russo, and Krysia Broda (2015). *The ILASP system for learning answer set programs*.

Li, Shen, Hengru Xu, and Zhengdong Lu (2018). "Generalize symbolic knowledge with neural rule engine". In: *arXiv preprint arXiv:1808.10326*.

Payani, Ali and Faramarz Fekri (2019). "Inductive Logic Programming via Differentiable Deep Neural Logic Networks". In: *CoRR* abs/1906.03523. arXiv: 1906.03523. URL: http://arxiv.org/abs/1906.03523.

Bibliography IV

Peng, Hao et al. (2018). "Rational Recurrences". In: *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*. Brussels, Belgium: Association for Computational Linguistics, pp. 1203–1214. DOI: 10.18653/v1/D18-1152. URL: https://www.aclweb.org/anthology/D18-1152.

Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin (2016a). ""Why Should I Trust You?": Explaining the Predictions of Any Classifier". In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Francisco, CA, USA, August 13-17, 2016*, pp. 1135–1144.

Ribeiro, Marco Túlio, Sameer Singh, and Carlos Guestrin (2016b). ""Why Should I Trust You?": Explaining the Predictions of Any Classifier". In: *CoRR* abs/1602.04938. arXiv: 1602.04938. URL: http://arxiv.org/abs/1602.04938.

Schuster, Sebastian et al. (2018). "Cross-lingual transfer learning for multilingual task oriented dialog". In: *arXiv preprint arXiv:1810.13327*.

Bibliography V

Schwartz, Roy, Sam Thomson, and Noah A. Smith (July 2018).
   "Bridging CNNs, RNNs, and Weighted Finite-State Machines". In:
   *Proceedings of the 56th Annual Meeting of the Association for
   Computational Linguistics (Volume 1: Long Papers)*. Melbourne,
   Australia: Association for Computational Linguistics, pp. 295–305.
   DOI: 10.18653/v1/P18-1028. URL:
   https://www.aclweb.org/anthology/P18-1028.

Suresh, Ananda Theertha et al. (2019). "Approximating probabilistic
   models as weighted finite automata". In: *CoRR* abs/1905.08701.
   arXiv: 1905.08701. URL: http://arxiv.org/abs/1905.08701.

Viterbi, Andrew (1967). "Error bounds for convolutional codes and an
   asymptotically optimum decoding algorithm". In: *IEEE transactions
   on Information Theory* 13.2, pp. 260–269.

Wan, Alvin et al. (2020). "NBDT: Neural-Backed Decision Trees". In:
   arXiv preprint arXiv:2004.00221.

Bibliography VI

Wang, Cheng and Mathias Niepert (2019). "State-Regularized
  Recurrent Neural Networks". In: ed. by Kamalika Chaudhuri and
  Ruslan Salakhutdinov. Vol. 97. Proceedings of Machine Learning
  Research. Long Beach, California, USA: PMLR, pp. 6596–6606. URL:
  http://proceedings.mlr.press/v97/wang19j.html.

Yin, Penghang et al. (2019). "Understanding straight-through estimator
  in training activation quantized neural nets". In: *arXiv preprint
  arXiv:1903.05662.*