

SoPa++: Leveraging explainability from hybridized RNN, CNN and weighted finite-state neural architectures

Atreya Shankar, shankar@uni-potsdam.de
Cognitive Systems: Language, Learning, and Reasoning (M.Sc.)

Thesis Progress Update
University of Potsdam, WiSe 20/21

February 3, 2021

Recap

- 1 Source code refactored to work with latest PyTorch and CUDA versions
- 2 Facebook NLU multi-class intent detection data set from [Schuster et al. \(2018\)](#)
- 3 Performance is important for explainability, because it does not make sense to explain a poorly performing model
- 4 Focus would be to develop SoPa++ as a neural model which can be easily decomposed into a *globally* explainable model

Semirings

Kuich and Salomaa 1986:

A semiring is a set \mathbb{K} along with two binary associative operations \oplus (addition) and \otimes (multiplication) and two identity elements: $\bar{0}$ for addition and $\bar{1}$ for multiplication. Semirings require that addition is commutative, multiplication distributes over addition, and that multiplication by $\bar{0}$ annihilates, i.e., $\bar{0} \otimes a = a \otimes \bar{0} = \bar{0}$

- Generic semiring notation: $\langle \mathbb{K}, \oplus, \otimes, \bar{0}, \bar{1} \rangle$
- Max-sum semiring: $\langle \mathbb{R} \cup \{-\infty\}, \max, +, -\infty, 0 \rangle$
- Max-product semiring: $\langle \mathbb{R} \cup \{-\infty\}, \max, \times, -\infty, 1 \rangle$
- Many semirings are possible, but we stick to *max* based semirings

Weighted Finite-State Automata

$$p(\mathbf{x}) = \boldsymbol{\pi} \otimes \left(\bigotimes_{i=1}^n \mathbf{T}(x_i) \right) \otimes \boldsymbol{\eta} \quad (1)$$

$$s(\mathbf{z}) = \bigoplus_{\mathbf{x} \in \Pi(\mathbf{z})} p(\mathbf{x}) \quad (2)$$

- Path $\mathbf{x} = \langle x_1, x_2, \dots, x_n \rangle$ derives some string $\mathbf{z} = \langle z_1, z_2, \dots, z_m \rangle$
- $\Pi(\mathbf{z})$ is the set of all possible paths that derive string \mathbf{z}
- $\boldsymbol{\pi}, \boldsymbol{\eta}$ are start and end state vectors
- \mathbf{T} is a transition matrix defining transition scores between states
- $p(\mathbf{x})$ refers to a score through a particular path
- $s(\mathbf{z})$ refers to an aggregate score through all possible paths
- Utilize Viterbi algorithm with *max* based semirings to compute $s(\mathbf{z})$
(Viterbi, 1967)

Finite transition types




Figure 1: Schematic for finite-state transitions; ϵ represents an epsilon-transition, “funny” and “magical” form self-loops and all other transitions are main transitions (Schwartz, Thomson, and Smith, 2018)

- Vanilla SoPa allowed 3 transitions types: main, epsilon and self-loops
- **SoPa++ uses only two transitions: namely main and wildcard transitions**
- Both consume one token and one state, which leads to pattern-string determinism and could make the *explainability* step easier

SoPa++: Lower Model




Figure 2: SoPa++ lower model schematic where pattern 1 is activated by the substring "in years" and pattern 2 is activated by the substring "funniest and most likeable book" (Schwartz, Thomson, and Smith, 2018)

- Vanilla SoPa places a MLP on top of the max-pooled states to retrieve classification results
- This contributes to a lack of explainability since MLPs are generally not easily explainable

SoPa++: Upper Model

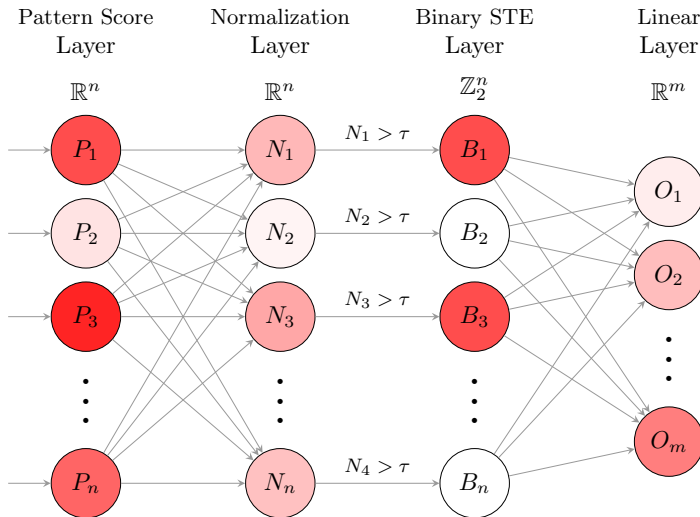


Figure 3: SoPa++ upper model schematic (combined lower-upper schematic under-construction); STE refers to a straight-through-estimator (Yin et al., 2019)

SoPa++: Explainable Model

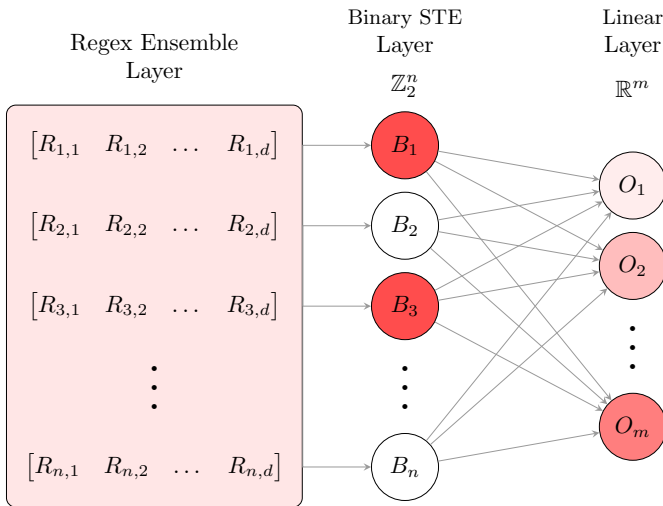


Figure 4: SoPa++ explainable variant derived from antecedent neural model

Training and Convergence



Figure 5: Excerpt from TensorBoard training logs for SoPa++ grid-search

- Model training procedure tends to converge fast
- Best model checkpoint(s) were reached within 1-6 epochs

Evaluation

Patterns	Test Precision	Test Recall	Test F_1
6-10_5-10_4-10_3-10	0.982	0.981	0.982
6-50_5-50_4-50_3-50	0.988	0.987	0.987

Table 1: Performance summary of various SoPa++ models; evaluation metrics are based on weighted averages

- Performance is competitive with that from literature and online benchmarks, which show F_1 scores ranging from 96-99% ([Schuster et al., 2018](#))
- In both the above cases, the max-sum semiring outperforms the max-product semiring
- Even though the SoPa++ model variant with higher number of patterns performs better, a **smaller model** might be more beneficial for explainability

Explainability: Peek into the Regex Ensemble

Pattern length	Sample regular expressions
3	{* forecast}, {* activate}, {* weather}
4	{activate morning *}, {my alarm *}, {* minute snooze}
5	{add * reminder for}, {the boiled eggs *}
6	{to wish * * happy}, {* add * * to}

Table 2: Sample regular expressions that activate patterns of specified lengths; * is used for aesthetic purposes to represent a wildcard instead of the appropriate `\w+` regular expression

- Patterns above were processed from a small subset of the training data; the full processing is still under development
- Patterns are indicative of important (generalized) n-grams in text
- Analyzing the weights of the linear layer can tell us how important each pattern is for the prediction of each class

Next steps

- 1 Find a clever way for clustering/merging regular expressions to keep the size of explainable model small
- 2 Analyze the performance of both SoPa++ and its explainable variant on the test set to measure how different their performances are
- 3 For cases where the explainable model does not perform similar to the SoPa++ neural model, propose local sample explanations as alternative

Bibliography I

- Arrieta, Alejandro Barredo et al. (2020). “Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI”. In: *Information Fusion* 58, pp. 82–115.
- Bocklisch, Tom et al. (2017). “Rasa: Open source language understanding and dialogue management”. In: *arXiv preprint arXiv:1712.05181*.
- Cybenko, George (1989). “Approximation by superpositions of a sigmoidal function”. In: *Mathematics of control, signals and systems* 2.4, pp. 303–314.
- Doran, Derek, Sarah Schulz, and Tarek R. Besold (2017). “What Does Explainable AI Really Mean? A New Conceptualization of Perspectives”. In: *CoRR* abs/1710.00794. arXiv: 1710.00794. URL: <http://arxiv.org/abs/1710.00794>.

Bibliography II

Evans, Richard and Edward Grefenstette (2018). “Learning explanatory rules from noisy data”. In: *Journal of Artificial Intelligence Research* 61, pp. 1–64.

Hornik, Kurt, Maxwell Stinchcombe, Halbert White, et al. (1989). “Multilayer feedforward networks are universal approximators.”. In: *Neural networks* 2.5, pp. 359–366.

Hou, Bo-Jian and Zhi-Hua Zhou (2018). “Learning with Interpretable Structure from RNN”. In: *CoRR* abs/1810.10708. arXiv: 1810.10708. URL: <http://arxiv.org/abs/1810.10708>.

Jiang, Chengyue et al. (2020). “Cold-Start and Interpretability: Turning Regular Expressions into Trainable Recurrent Neural Networks”. In:

Kepner, Jeremy et al. (2018). “Sparse deep neural network exact solutions”. In: *2018 IEEE High Performance extreme Computing Conference (HPEC)*. IEEE, pp. 1–8.

Bibliography III

Kuich, Werner and Arto Salomaa (1986). “Linear Algebra”. In: *Semirings, automata, languages*. Springer, pp. 5–103.

Law, Mark, Alessandra Russo, and Krysia Broda (2015). *The ILASP system for learning answer set programs*.

Li, Shen, Hengru Xu, and Zhengdong Lu (2018). “Generalize symbolic knowledge with neural rule engine”. In: *arXiv preprint arXiv:1808.10326*.

Payani, Ali and Faramarz Fekri (2019). “Inductive Logic Programming via Differentiable Deep Neural Logic Networks”. In: *CoRR* abs/1906.03523. arXiv: 1906.03523. URL: <http://arxiv.org/abs/1906.03523>.

Bibliography IV

Peng, Hao et al. (2018). “Rational Recurrences”. In: *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*. Brussels, Belgium: Association for Computational Linguistics, pp. 1203–1214. DOI: 10.18653/v1/D18-1152. URL: <https://www.aclweb.org/anthology/D18-1152>.

Ribeiro, Marco Tulio, Sameer Singh, and Carlos Guestrin (2016a). ““Why Should I Trust You?”: Explaining the Predictions of Any Classifier”. In: *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Francisco, CA, USA, August 13-17, 2016*, pp. 1135–1144.

Ribeiro, Marco Túlio, Sameer Singh, and Carlos Guestrin (2016b). ““Why Should I Trust You?”: Explaining the Predictions of Any Classifier”. In: *CoRR* abs/1602.04938. arXiv: 1602.04938. URL: <http://arxiv.org/abs/1602.04938>.

Schuster, Sebastian et al. (2018). “Cross-lingual transfer learning for multilingual task oriented dialog”. In: *arXiv preprint arXiv:1810.13327*.

Bibliography V

Schwartz, Roy, Sam Thomson, and Noah A. Smith (July 2018).

“Bridging CNNs, RNNs, and Weighted Finite-State Machines”. In: *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Melbourne, Australia: Association for Computational Linguistics, pp. 295–305. DOI: 10.18653/v1/P18-1028. URL: <https://www.aclweb.org/anthology/P18-1028>.

Suresh, Ananda Theertha et al. (2019). “Approximating probabilistic models as weighted finite automata”. In: *CoRR* abs/1905.08701. arXiv: 1905.08701. URL: <http://arxiv.org/abs/1905.08701>.

Viterbi, Andrew (1967). “Error bounds for convolutional codes and an asymptotically optimum decoding algorithm”. In: *IEEE transactions on Information Theory* 13.2, pp. 260–269.

Wan, Alvin et al. (2020). “NBDT: Neural-Backed Decision Trees”. In: *arXiv preprint arXiv:2004.00221*.

Bibliography VI

- Wang, Cheng and Mathias Niepert (2019). “State-Regularized Recurrent Neural Networks”. In: ed. by Kamalika Chaudhuri and Ruslan Salakhutdinov. Vol. 97. Proceedings of Machine Learning Research. Long Beach, California, USA: PMLR, pp. 6596–6606. URL: <http://proceedings.mlr.press/v97/wang19j.html>.
- Yin, Penghang et al. (2019). “Understanding straight-through estimator in training activation quantized neural nets”. In: *arXiv preprint arXiv:1903.05662*.